

6.034: Introduction to Artificial Intelligence

Adversarial search & games

Lecture 5 Handout

Robert C. Berwick
September 14, 2020

Menu for today

Games: Adversarial search


- ☐ Intelligence & Chess
- ☐ Ways to Play Chess
- ☐ How to search: Minimax algorithm
- ☐ Improving search: Alpha-beta pruning
- ☐ Progressive Deepening
- ☐ Reflections on contemporary chess computers & Gold Star ideas

Games: Adversarial Search

How could a computer play chess?

1. Human-like
 2. If then: make most plausible move
 - * 3. Look ahead & evaluate
 - evaluate moves by points, etc.
- $S = g(f_1, f_2, \dots, f_n)$
score
4. British museum algorithm
 - though ^{not} that productive

Vocabulary for game trees




Could we use cloud computing to evaluate all possible chess moves in the same way?


10^{120} moves

$$\frac{\frac{10^{80}}{\pi \times 10^7}}{\frac{10^9}{10^{10}}} \frac{\text{atoms in universe}}{\text{seconds/year}} \frac{\text{nanoseconds/second}}{\text{years since Big Bang}}$$

$\approx 10^{26}$ nanoseconds x 10^{80} atoms
 $\approx 10^{106}$ ops, each atom operating
@ nanosecond speeds
Not enough compute power for
 10^{120} moves!

Minimax algorithm for searching game tree to find optimal move sequence





Note that the best minimax path is preserved under any transform of the leaf scores that preserves their rank ordering



Games

17:13:35 EDT 02-Aug-2020

- Goal trees
- Search
 - Basic search
 - Optimal search
 - Games**
 - Monte Carlo
- Constraint satisfaction
- Biological mimetics
- Learning
- Neural nets
- Bayes nets
- Miscellaneous

Start New tree Hide choices

Theoretical minimum / actual / maximum: 7 / 16 / 16
Minimax only


MAX

MIN

MAX

MIN

Static
Values



Example

- ☒ Classroom example
- ☐ Quiz example
- ☐ Good news example

Arrangement

- ☐ Random numbers
- ☐ Increasing left to right
- ☐ Decreasing left to right

Search type

- ☒ Mini-max only
- ☐ Mini-max with alpha-beta

Shape


- ☒ 2 x 4
- ☐ 4 x 2
- ☐ 2 x 6
- ☐ 6 x 2

Delay



Minimax with alpha-beta (α - β) pruning

don't have to evaluate all static values of leaf nodes



α value = MAX is guaranteed to gain at least this much, or $\geq \alpha$; a floor on MAX's gain.

β value = MIN is guaranteed to lose at most this much, or $\leq \beta$; a ceiling on MIN's loss.


Note: $[\alpha < \beta]$, or floor < ceiling. Why?

Q: Alpha-Beta pruning value = Minimax value w/o pruning?


Ans: 2

prune if $\alpha \geq \beta$


A deeper game tree illustrates how much pruning alpha-beta can do ($b = 2$, $d = 3$)



Second half of tree: cutoff whenever $\alpha \geq \beta$



A deeper game tree



6 out of 16 static evaluations are not made



Dead horse principle (principle AWP)



6034

13:22:46 EDT 02-Aug-2020

Start New tree Hide choices

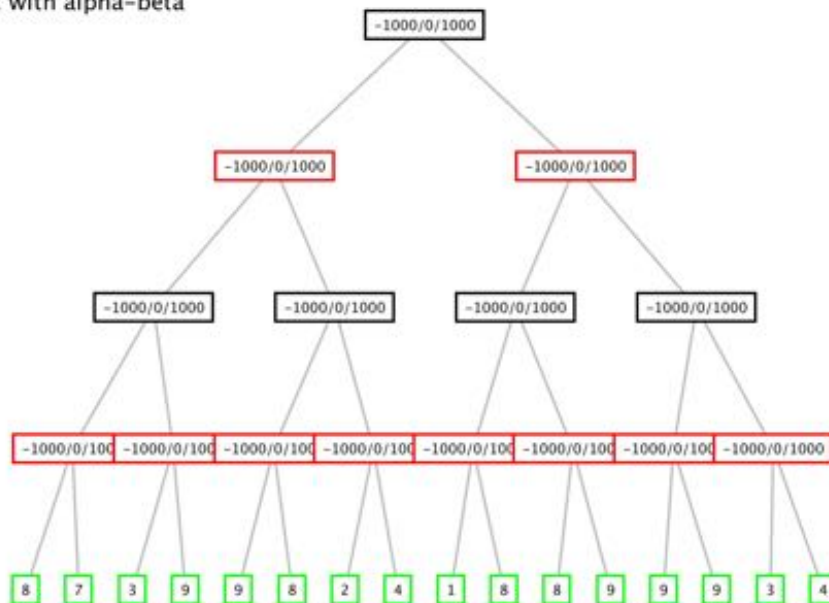
Theoretical minimum / actual / maximum: 7 / 0 / 16
Minimax with alpha-beta

MAX

MIN

MAX

MIN



Example

- Classroom example
- Quiz example
- Good news example

Arrangement

- ☐ Random numbers
- ☐ Increasing left to right
- ☐ Decreasing left to right

- Search type

- ☐ Mini-max only
- ☒ Mini-max with alpha-beta

Shape

- ☒ 2 x 4
☐ 4 x 2
☐ 2 x 6
☐ 6 x 2


Delay



Ordering of static values can greatly affect alpha-beta pruning

- If the most favorable successor nodes for both MAX and MIN are on the LEFT so we explore them FIRST, then this leads to maximal pruning
- If the most favorable successor nodes for both MAX and MIN are on the RIGHT so they are explored LAST, then there is less pruning, possibly none at all
- Maximal pruning (in terms of branching factor b and tree depth d is approx: $2 b^{d/2}$), so can search down 2x as far using α - β search in this “good news” optimal case, compared to std minimax

Prune whenever $\alpha \geq \beta$

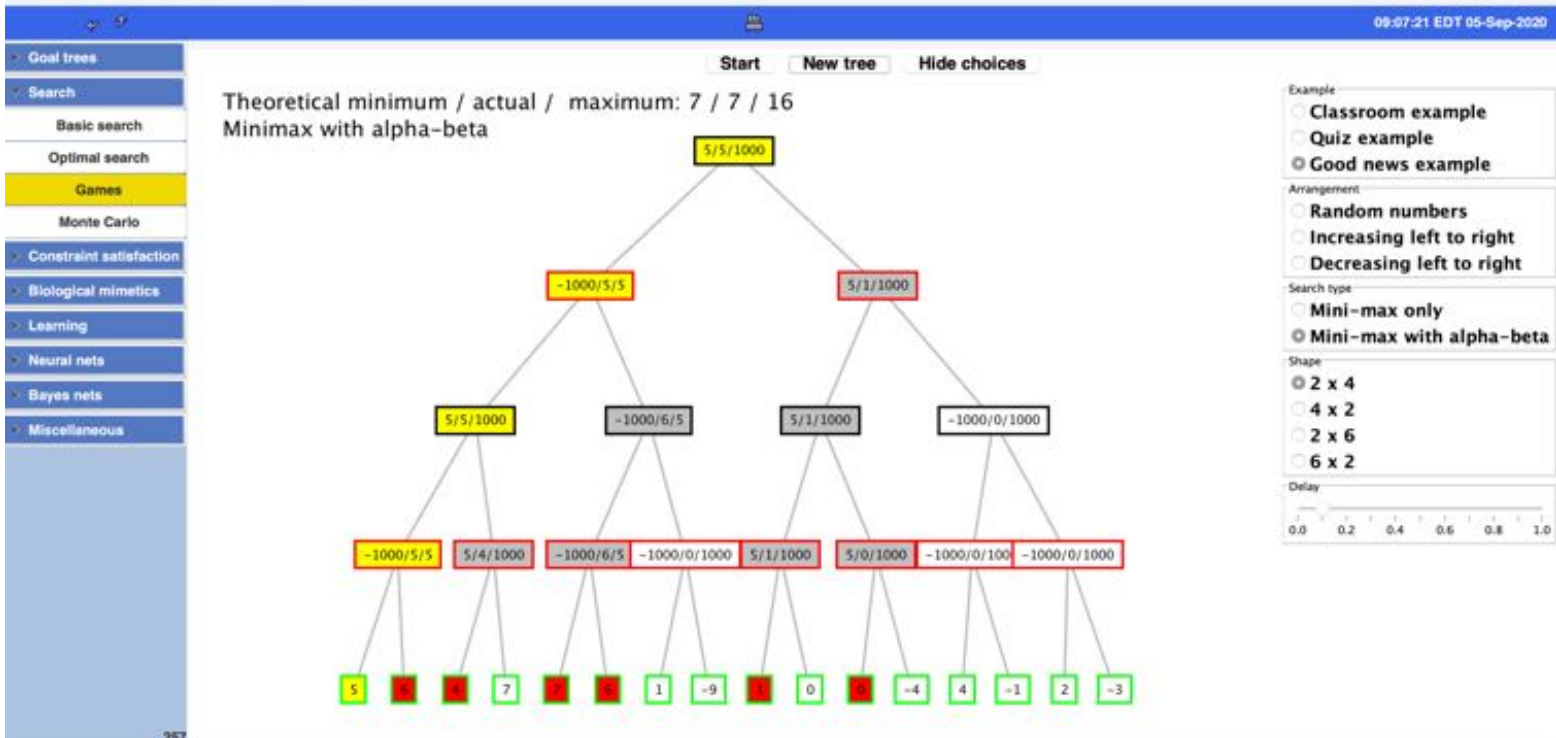


What would happen to pruning if nodes 3, 9 reversed? 3 Not pruned!

Optimal (good news) α - β game tree pruning



Games



$\approx 2b^{d/2}$ \approx maximum savings using α - β if
optimally ordered game tree

But suppose we run out of compute time?
The branching factor depends on the game &
board state – won't know for sure how deep
we can go in 1 minute....

Suppose we are playing blitz chess?

Will we always have a (good) move at hand?

Let's combine this thought with the idea of
optimally ordering the node evaluations...




Anytime algorithm




Martial arts principle

How not to run out of time



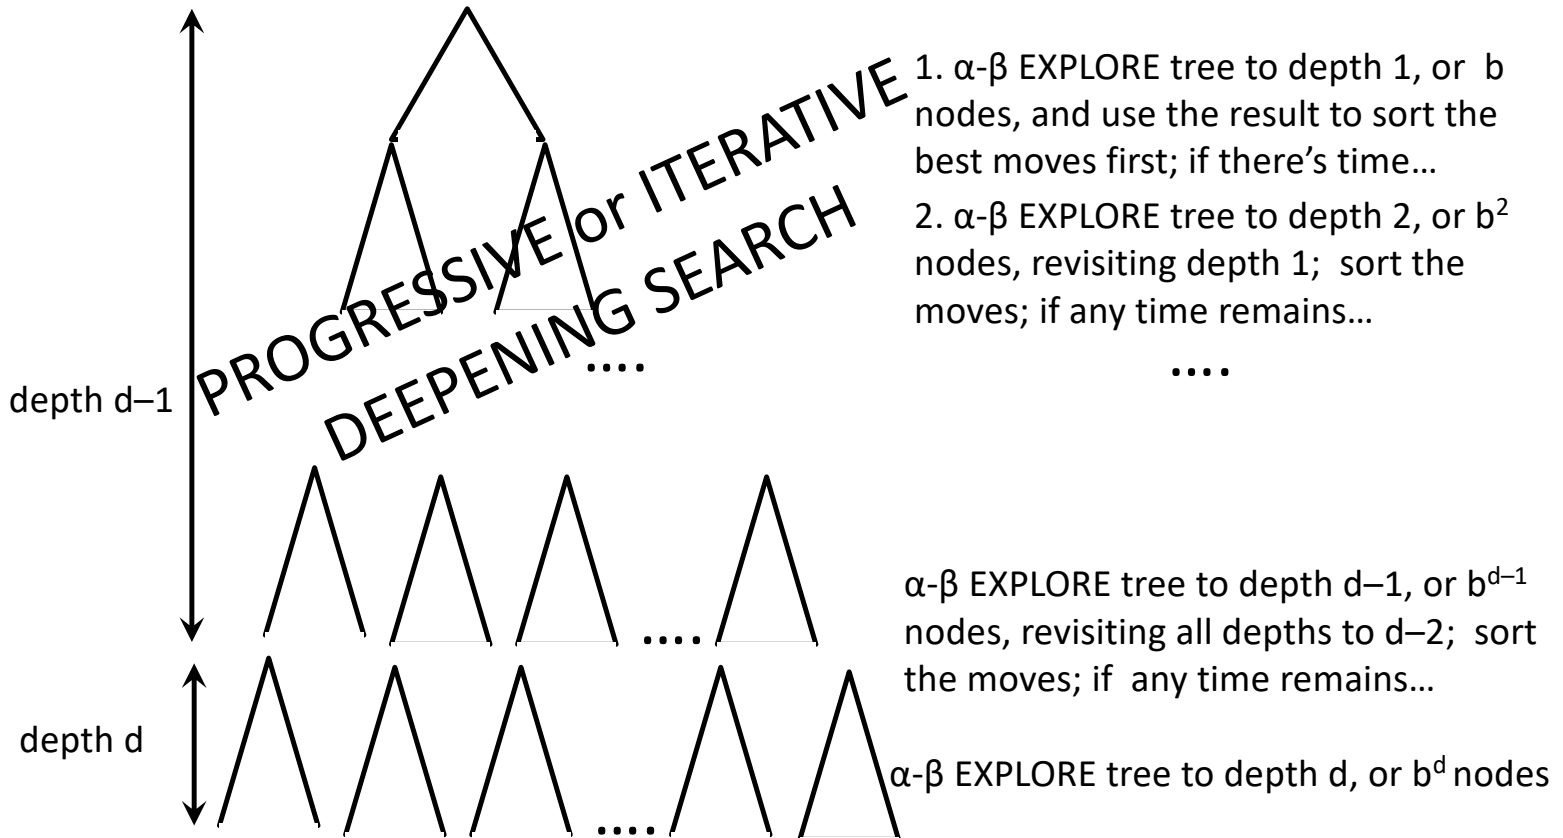
....



So, we can have a move in hand at level d-1 for only $1/b^{\text{th}}$ of computation cost required to go all the way down to level d


We can do this recursively to great benefit...

How not to run out of time (ignoring α - β for now)



Computational cost of progressive deepening?

How many total nodes S in game tree to depth $d-1$?



Progressive (Iterative) deepening search

- Optimal DFS method (redundant use of space, but doesn't lose on arbitrary depth d trees)
- Earlier searches tend to improve the commonly used heuristics, so that a more accurate estimate of the score of various nodes at the final depth search can occur
- Because early iterations use small values for d they execute extremely quickly. This allows the algorithm to supply early indications of the result almost immediately, followed by refinements as d increases

Gold star ideas today

- ★ Dead horse principle, aka “AWP”: α - β search
- ★ Martial arts principle – use adversary’s strength against them: progressive deepening
- ★ Anytime algorithms: progressive deepening
- ★ Simple \neq Trivial: sometimes, bulldozers work